

Predicting Breast Cancer using logistic regression classification

[Code ▾](#)[Hide](#)

```
# Importing the dataset
dataset = read.csv('breast_cancer_prediction.csv')
dataset = data.frame(dataset)
# Encoding the target feature as factor
dataset$Classification = ifelse(dataset$Classification==1, 0,1)
head(dataset)
```

...	BMI	Glucose	Insulin	HOMA	Leptin	Adiponectin	Resistin	MCP.1
<int>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1 48	23.50000	70	2.707	0.4674087	8.8071	9.702400	7.99585	417.114
2 83	20.69049	92	3.115	0.7068973	8.8438	5.429285	4.06405	468.786
3 82	23.12467	91	4.498	1.0096511	17.9393	22.432040	9.27715	554.697
4 68	21.36752	77	3.226	0.6127249	9.8827	7.169560	12.76600	928.220
5 86	21.11111	92	3.549	0.8053864	6.6994	4.819240	10.57635	773.920
6 49	22.85446	92	3.226	0.7320869	6.8317	13.679750	10.31760	530.410

6 rows | 1-10 of 10 columns

[Hide](#)

```
# Splitting the dataset into the Training set and Test set
# install.packages('caTools')
library(caTools)
set.seed(123)
split = sample.split(dataset$Classification, SplitRatio = 0.8)
training_set = subset(dataset, split == TRUE)
test_set = subset(dataset, split == FALSE)
# Feature Scaling
out_index = which(colnames(dataset)=="Classification")
training_set[-out_index] = scale(training_set[-out_index])
test_set[-out_index] = scale(test_set[-out_index])
# Fitting Logistic Regression to the Training set
classifier = glm(formula = Classification ~ .,
                  family = binomial,
                  data = training_set)
summary(classifier)
```

```
Call:
glm(formula = Classification ~ ., family = binomial, data = training_set)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.1095	-0.7409	0.1632	0.6544	2.2083

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.5019	0.3670	1.368	0.1714
Age	-0.3806	0.2888	-1.318	0.1876
BMI	-0.6685	0.3778	-1.769	0.0768 .
Glucose	2.0568	0.8016	2.566	0.0103 *
Insulin	1.9181	3.1383	0.611	0.5411
HOMA	-1.5071	4.2230	-0.357	0.7212
Leptin	-0.4548	0.3811	-1.193	0.2327
Adiponectin	0.1215	0.3107	0.391	0.6958
Resistin	0.8091	0.3739	2.164	0.0305 *
MCP.1	0.2876	0.3165	0.909	0.3635

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


(Dispersion parameter for binomial family taken to be 1)

Null deviance: 128.053 on 92 degrees of freedom
Residual deviance: 82.904 on 83 degrees of freedom
AIC: 102.9

Number of Fisher Scoring iterations: 7

Hide

```
# Interesting results from the model, blood glucose levels and resistin levels are statistically
# significant at p<0.05 in predicting breast cancer in trainign data.
# Predicting the Test set results
prob_pred = predict(classifier, type = 'response', newdata = test_set[-out_index])
y_pred = ifelse(prob_pred > 0.5, 1, 0)
# ROC Curve
library(pROC)
preds=predict(classifier,test_set[-out_index], type="response")
ro <- roc(test_set[,out_index] ~ preds)
plot(ro)
```



Hide

```
auc(ro)
```

Area under the curve: 0.6231

Hide

```
# Evaluating Model Accuracy on test data set using
# Confusion Matrix
cm = table(test_set[, out_index], y_pred > 0.5)
print(cm)
```

```
FALSE TRUE
0      8   2
1      7   6
```

Hide

```
Model_Accuracy=(cm[1,1]+cm[2,2])/(cm[1,1]+cm[1,2]+cm[2,1]+cm[2,2])
print("Assuming 50% as cutoff, Model Accuracy is")
```

```
[1] "Assuming 50% as cutoff, Model Accuracy is"
```

Hide

```
print(Model_Accuracy)
```

[1] 0.6086957